Why AI Answers Have Different Costs

Tokens, computation, waiting and reuse

SCAN / LEARN

THE MAIN IDEA

AI costs depend on the amount of work performed—not merely on the number of questions. Good routing avoids needless work.

FOLLOW THE SYSTEM

01

INPUT SIZE

More input tokens require processing.

›

02

MODEL CHOICE

Different models have different speed and cost.

›

03

OUTPUT & TIME

Longer generation uses more computation and time.

›

04

REUSE

Safe caching can avoid some repeated calculations.

TOKENS + MODEL + TIME − SAFE REUSE = COST

COMMON MISCONCEPTION

The fastest or cheapest response is always the right choice.

MORE ACCURATE

Quality, correctness, privacy and delay matter alongside cost.

THREE QUESTIONS TO KEEP

PROBLEM: How do we spend compute intelligently?

FAILURE: Waste and long waits can overwhelm a service.

IN REAL LIFE: Chat systems, agents, search, support.

SOLO DISCOVERY

Try alone: compare the effort of answering “Hi” with summarizing a long report.