Can We Trust This AI System?

Evaluation is more than a good demo

SCAN / LEARN

THE MAIN IDEA

A strong system succeeds on representative tasks, repeatedly and safely—not just once on a polished example.

FOLLOW THE SYSTEM

01 DEFINE THE TASK

Agree on what counts as a correct result.

02 BUILD TESTS

Include typical, difficult and unexpected cases.

03 MEASURE

Track accuracy, failures, time and cost.

04 RETEST

Repeat after updates and inspect real-world drift.

REQUIREMENTS → TESTS → EVIDENCE → IMPROVEMENT

COMMON MISCONCEPTION

One impressive answer proves a system is reliable.

MORE ACCURATE

A repeatable test set and honest failure reporting are stronger evidence.

THREE QUESTIONS TO KEEP

PROBLEM: How do we know a tool works well?

FAILURE: A demo hides fragile or unsafe behavior.

IN REAL LIFE: Software releases, medical support, education.

SOLO DISCOVERY

Try alone: write three questions a study assistant must answer correctly, plus one trick question.